

Week 11 Report

CNN

37237256 薛丁铭

The code of this week is as shown in Figure 1. I tried to add 4 convolutional layers to the network but it showed error messages. I think the limit should be 3 because the network has to maintain its shape after pooling. The accuracy increased with epochs. The results are as shown in Figure 2.

```
school > Introduction to Machine Learning > Week11_exercise > CNN.py > ...
1 import matplotlib.pyplot as plt
2 import tensorflow as tf
3 (x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()
4 print(x_train.shape)
5
6 # plt.imshow(x_train[1])
7 # plt.show()
8
9 x_train, x_test = x_train/255.0, x_test/255.0
10
11 ## Regular NN
12 # model_NN = tf.keras.models.Sequential([
13 #     tf.keras.layers.Flatten(input_shape=(32,32,3)),
14 #     tf.keras.layers.Dense(128, activation='sigmoid'),
15 #     tf.keras.layers.Dense(128, activation='sigmoid'),
16 #     tf.keras.layers.Dense(128, activation='relu'),
17 #     tf.keras.layers.Dense(64, activation='relu'),
18 #     tf.keras.layers.Dense(64, activation='relu'),
19 #     tf.keras.layers.Dense(32, activation='relu'),
20 #     tf.keras.layers.Dense(100, activation='softmax')
21 # ])
22
23 # model_NN.compile(
24 #     optimizer = 'adam',
25 #     loss = 'sparse_categorical_crossentropy',
26 #     metrics = ['accuracy']
27 # )
28
29 # model_NN.fit(x_train, y_train, epochs = 10, validation_data = (x_test, y_test), batch_size = 128)
30
```

Figure 1.a Code

```
## CNN
model_CNN = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(32, (3,3), activation='relu', input_shape=(32,32,3), strides = (1,1)),
    tf.keras.layers.AveragePooling2D((2,2)),

    tf.keras.layers.Conv2D(64, (3,3), activation='relu', strides = (1,1)),
    tf.keras.layers.AveragePooling2D((2,2)),

    tf.keras.layers.Conv2D(128, (3,3), activation='relu', strides = (1,1)),
    tf.keras.layers.MaxPooling2D((2,2)),

    # tf.keras.layers.Conv2D(128, (3,3), activation='relu', strides = (1,1)),
    # tf.keras.layers.MaxPooling2D((2,2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(32, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])

model_CNN.summary()
model_CNN.compile(
    optimizer = 'adam',
    loss = 'sparse_categorical_crossentropy',
    metrics = ['accuracy']
)

# model_CNN.fit(x_train, y_train, epochs = 5, validation_data = (x_test, y_test), batch_size = 128)

history_CNN = model_CNN.fit(x_train, y_train, epochs = 20, validation_data = (x_test, y_test), batch_size = 128)
plt.plot(history_CNN.history['accuracy'], label = 'accuracy')
plt.plot(history_CNN.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0.5,1])
plt.legend(loc = 'lower right')
plt.show()
```

Figure 1.b Code

```
(50000, 32, 32, 3)
2024-01-06 18:43:05.858770: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized with one
n performance-critical operations: AVX AVX2
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.
Model: "sequential"

-----
Layer (type)                Output Shape                Param #
-----
conv2d (Conv2D)              (None, 30, 30, 32)         896
average_pooling2d (AverageP (None, 15, 15, 32)         0
ooling2D)
conv2d_1 (Conv2D)            (None, 13, 13, 64)         18496
average_pooling2d_1 (Averag (None, 6, 6, 64)           0
ePooling2D)
conv2d_2 (Conv2D)            (None, 4, 4, 128)          73856
max_pooling2d (MaxPooling2D (None, 2, 2, 128)          0
)
flatten (Flatten)            (None, 512)                 0
dense (Dense)                 (None, 128)                 65664
dense_1 (Dense)               (None, 128)                 16512
dense_2 (Dense)               (None, 64)                  8256
dense_3 (Dense)               (None, 32)                  2080
dense_4 (Dense)               (None, 10)                  330
-----
Total params: 186,090
```

Figure 2.a Results

```
-----
Total params: 186,090
Trainable params: 186,090
Non-trainable params: 0
-----
Epoch 1/20
391/391 [=====] - 17s 42ms/step - loss: 1.8007 - accuracy: 0.3252 - val_loss: 1.5578 - val_accuracy: 0.4271
Epoch 2/20
391/391 [=====] - 35s 90ms/step - loss: 1.4155 - accuracy: 0.4819 - val_loss: 1.3469 - val_accuracy: 0.5093
Epoch 3/20
391/391 [=====] - 42s 108ms/step - loss: 1.2447 - accuracy: 0.5487 - val_loss: 1.2074 - val_accuracy: 0.5623
Epoch 4/20
391/391 [=====] - 35s 88ms/step - loss: 1.1202 - accuracy: 0.6006 - val_loss: 1.1064 - val_accuracy: 0.6047
Epoch 5/20
391/391 [=====] - 35s 90ms/step - loss: 1.0269 - accuracy: 0.6362 - val_loss: 1.0211 - val_accuracy: 0.6324
Epoch 6/20
391/391 [=====] - 34s 86ms/step - loss: 0.9520 - accuracy: 0.6606 - val_loss: 1.0410 - val_accuracy: 0.6372
Epoch 7/20
391/391 [=====] - 41s 105ms/step - loss: 0.8789 - accuracy: 0.6888 - val_loss: 0.9629 - val_accuracy: 0.6586
Epoch 8/20
391/391 [=====] - 46s 118ms/step - loss: 0.8128 - accuracy: 0.7143 - val_loss: 0.9405 - val_accuracy: 0.6698
Epoch 9/20
391/391 [=====] - 35s 89ms/step - loss: 0.7666 - accuracy: 0.7301 - val_loss: 0.9180 - val_accuracy: 0.6794
Epoch 10/20
391/391 [=====] - 34s 87ms/step - loss: 0.7169 - accuracy: 0.7474 - val_loss: 0.9026 - val_accuracy: 0.6891
Epoch 11/20
391/391 [=====] - 33s 85ms/step - loss: 0.6830 - accuracy: 0.7601 - val_loss: 0.8952 - val_accuracy: 0.6943
Epoch 12/20
391/391 [=====] - 33s 85ms/step - loss: 0.6299 - accuracy: 0.7800 - val_loss: 0.9290 - val_accuracy: 0.6923
Epoch 13/20
391/391 [=====] - 33s 84ms/step - loss: 0.5957 - accuracy: 0.7917 - val_loss: 0.8889 - val_accuracy: 0.7039
Epoch 14/20
391/391 [=====] - 34s 88ms/step - loss: 0.5612 - accuracy: 0.8027 - val_loss: 0.9069 - val_accuracy: 0.7073
Epoch 15/20
391/391 [=====] - 39s 99ms/step - loss: 0.5347 - accuracy: 0.8138 - val_loss: 0.8919 - val_accuracy: 0.7089
Epoch 16/20
391/391 [=====] - 39s 101ms/step - loss: 0.4969 - accuracy: 0.8268 - val_loss: 0.9487 - val_accuracy: 0.7058
Epoch 17/20
391/391 [=====] - 33s 85ms/step - loss: 0.4712 - accuracy: 0.8347 - val_loss: 0.9298 - val_accuracy: 0.7103
Epoch 18/20
391/391 [=====] - 33s 84ms/step - loss: 0.4428 - accuracy: 0.8444 - val_loss: 0.9074 - val_accuracy: 0.7123
Epoch 19/20
391/391 [=====] - 33s 83ms/step - loss: 0.4176 - accuracy: 0.8546 - val_loss: 0.9606 - val_accuracy: 0.7129
Epoch 20/20
391/391 [=====] - 33s 85ms/step - loss: 0.3890 - accuracy: 0.8632 - val_loss: 0.9916 - val_accuracy: 0.7109
||
```

Figure 2.b Results

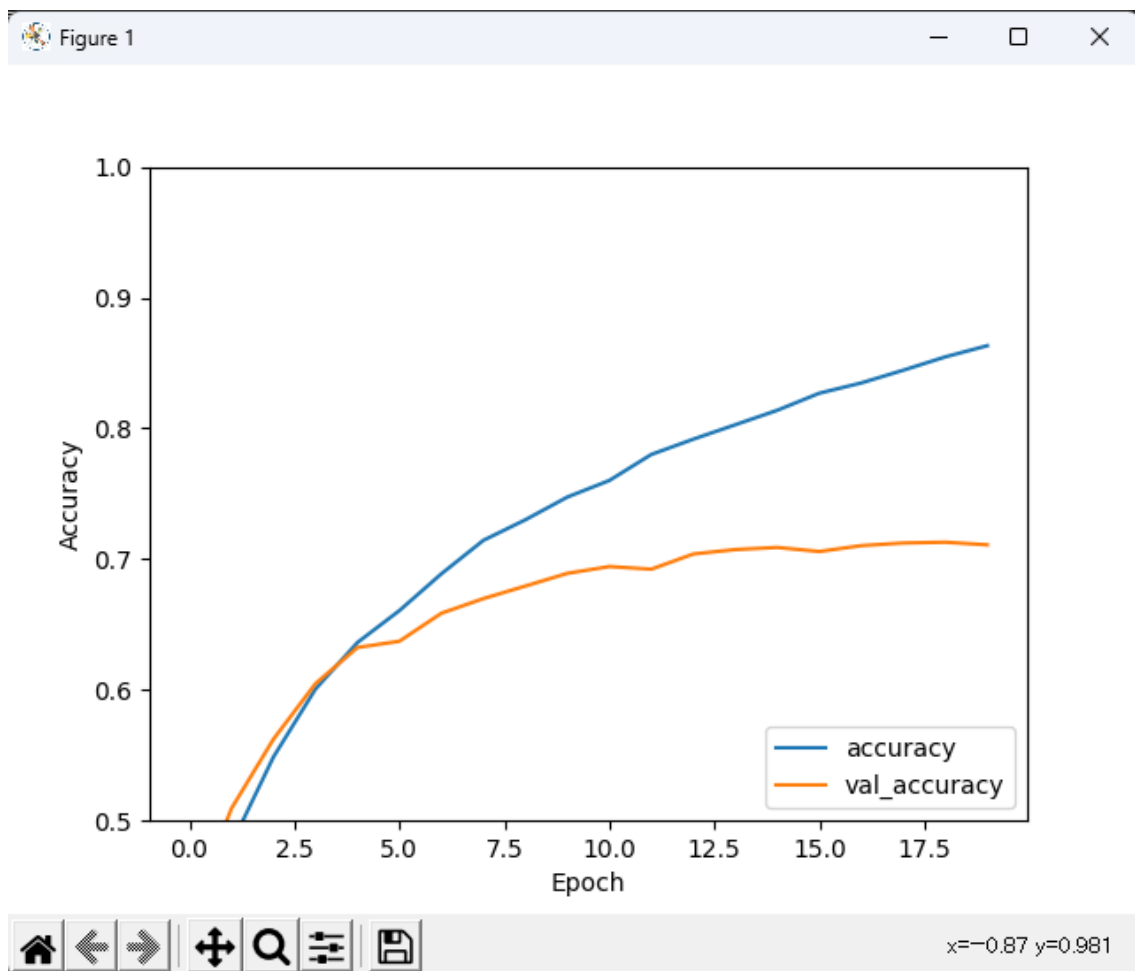


Figure 2.c Results